Gov 2001: Problem Set 4

Due Wednesday, March 2 by 6pm

Problem 1: Maximizing a Poisson Likelihood

Suppose $Y_i \sim Poisson(\lambda)$. You take two draws from the distribution of Y_i, y_1 and y_2.

1.A) Write out the expression for the likelihood function, $L(\lambda \mid y_1, y_2)$

1.B) Write out the expression for the log-likelihood function $\ell(\lambda) = \log [L(\lambda \mid y_1, y_2)]$

1.C) Write out the expression for the first derivative of the log-likelihood function $\ell'(\lambda) = \frac{\partial}{\partial \lambda} (\log [L(\lambda \mid y_1, y_2)])$

1.D) Write out the expression for the second derivative of the log-likelihood function $\ell''(\lambda) = \frac{\partial^2}{\partial \lambda^2} (\log [L(\lambda \mid y_1, y_2)])$

1.E) Find the MLE for λ by setting $\ell'(\lambda) = 0$ and checking that $\ell''(\lambda) < 0$.

1.F) Suppose $y_1 = 2$ and $y_2 = 4$. What is the value of your MLE, using the analytical solution you provided in part E?

1.G) Before we begin using `optim` to maximize likelihood and log-likelihood functions, we'll use this simple case to program our own numerical optimization. Recall from section that one common approach to numerical optimization is the Newton-Raphson method, where:

$$\lambda^{(i+1)} = \lambda^i - \frac{\ell'(\lambda^i)}{\ell''(\lambda^i)}$$

For an arbitrary starting value in the support of λ, $\lambda^{(0)}$, the expression above can be interated until resulting values of $\lambda^{(i+1)}$ converge to our MLE. Write out the update formula for the Newton Raphson algorithm that would calculate our MLE for λ.

1.H) Suppose again that $y_1 = 2$ and $y_2 = 4$. In R, use the update formula provided in part F to calculate your MLE numerically. Set $\lambda^{(0)}$ to 1.0. Does the algorithm converge to your answer in part F?

1.I) Finally, use the use `optim` with the following arguments `control(fnscale = -1)`, `method = "BFGS"`, `hessian = TRUE`. You can continue to assume that $y_1 = 2$ and $y_2 = 4$. Do your results match your results from F? Comment on the relationship between your hand-coded algorithm and this approach.

Hint: you'll need to write a function for your log-likelihood function from part B.

Problem 2: Maximizing a Normal Likelihood

$Y_i \sim N(\beta, 1)$ where β is constant across all Y_i. You receive the following data which consists of three realizations of Y: $(y_1 = -1, y_2 = 0, y_3 = 1)$. Assume all Y are independent of one another.

2.A) Find an expression for the log-likelihood of β conditional on your data. Remove any constants that will not affect the maximization of this function. Show your derivation.

2.B) Find the maximum likelihood estimate of β analytically.

2.C) Confirm that this is a maximum by checking the second derivative.

2.D) $-\frac{\partial^2 \ln L(\beta|\mathbf{y})}{\partial \beta^2}$, i.e. the negation of the 2nd derivative you calculated for 1.C, is an estimate of what is sometimes called 'the information'. Using the concept of the likelihood ratio (which on the log scale is a difference in log-likelihoods), explain in 3 sentences or less why this quantity might be a useful summary of the precision of the maximum likelihood estimate of a parameter.

2.E) Create a function in R for the log-likelihood from part A.

2.F) Confirm your answer from 1.B by maximizing your function from 1.E using `optim` with the following arguments `control(fnscale = -1)`, `method = "BFGS"`, `hessian = TRUE`. The hessian matrix is the matrix of second derivatives of a function, and note that in this case, it is 1 by 1 because our function has only one argument. Note that `optim` also provides the value of the likelihood at its maximum as `value`.

2.G) Use the Taylor Series expansion below to approximate the log-likelihood for β in the vicinity of its the maximum likelihood estimate.

Given $f(a)$, $f'(a)$, and $f''(a)$ where a is some value in the domain of f then

$$f(x) \approx f(a) + f'(a)(x - a) + \frac{f''(a)}{2}(x - a)^2.$$

Hint: for your problem f is the log-likelihood function and a is $\hat{\beta}_{MLE}$.

2.H) What do you think of this approximation for $\ln L(\beta|\mathbf{y})$? Based on your answer, would you consider your answer from part D a statement of the exact curvature or the approximate curvature of the likelihood around the maximum?

Problem 3

For this problem you will work with presidential election data. To begin, download the presidential.Rdata data file and load it into R using the following line of code:

```
load("presidential.Rdata")
```

These data, which are stored in matrix format, contain the dependent variable, Dvote, which is the fraction that the Democratic presidential candidate won in each state. The rest of the covariates include the following pieces of information for each state and year:

- year - Year of the race (from 1948 to 1992)

- n1 - Support for Dem candidate in Sept nationwide poll

- n2 - Presidential approval in July nationwide poll * Inc

- n3 - Presidential approval in July nationwide poll * Presinc

- n4 - 2nd Quarter GDP nationwide growth

- s1 - Dem share in the state the last election

- s2 - Dem share of state vote two elections ago

- s3 - Home state indicator variable

- s4 - Home state of VP candidate indicator

- s5 - Democratic minority in state legislature

- s6 - State economic growth in past year * Inc

- s7 - Measure of state ideology

- s8 - State ideological compatibility with candidates

- r1 - South indicator variable (i.e., the state is in the South)

- r2 - South in 1964*-1

- r3 - Deep South in 1964*-1

- r4 - New England in 1964

- r5 - North Central in 1972

- r6 - West in 1976*-1

Note that all variables are constructed so that an increase in the variable would be expected to increase Democratic vote share. Also, for the following problems, you may assume that (1) the observations are independent across states i and years t, and (2) the democratic vote share is normally distributed.

You may assume that σ^2 is constant from year to year and across states. That is, $Y_{it} \sim N(\mu_{it}, \sigma^2)$.

3.A) Parameterize the systematic component as $\mu_{it} = X_{it}\beta$. Using this parameterization, derive the log likelihood, $lnL(\beta, \sigma|y)$. Show the steps of your derivation.

3.B) Write a function in `R` that calculates this log likelihood. Your function should take as its inputs proposed parameter values for β, a vector for y, and a matrix for X.

3.C) Check that your function works, and then use the `optim` function to maximize the log likelihood for all covariates in the dataset. (Hint: For this problem, you'll want to give your `optim` function starting values of zero and use the method `BFGS`.) Report the MLE coefficient estimates for the Intercept, n_1, n_2, s_2, and r_3.
Be sure that you constrain your parameter for σ^2 to be positive by reparameterizing it in your function as e^{σ^2}.

Note: your model should use all covariates in the dataset, but you will only report a subset of them here. You can use the `lm()` function to check your work.

3.D) Extract and save the variance-covariance matrix from your optim output, and use this to calculate the standard errors associated with your coefficient estimates from Part C. Report the standard errors for the Intercept and n_1.

Problem 4: Likelihood Ratio Tests

A referee is challenging your model specification from Problem 3. He claims that your model would be a lot simpler and elegant if you just didn't include the regional variables: `r1`, `r2`, `r3`, `r4`, `r5`, and `r6` (conveniently coded in your dataset from problem 3!).

4.A) Specify the restricted and unrestricted models associated with the referee's claims.

4.B) Using `R` (and your function from Problem 3), calculate the likelihood ratio test statistic as a function of the two likelihoods. Report this value.

4.C) Calculate the probability of seeing this likelihood ratio test statistic under the null hypothesis that the restricted model is true. Do you accept the referee's suggestion of getting rid of the regional variables?

R code

Please submit all your code for this assignment as a .R file. Your code should be clean, commented, and executable without error.